

# Adversarial Search

Chapter 6

Section 1 – 4

# Warm Up

- Let's play some games!

# Outline

- Optimal decisions
- Imperfect, real-time decisions
- $\alpha$ - $\beta$  pruning


# Games vs. search problems

- "Unpredictable" opponent → specifying a move for every possible opponent reply
- Time limits → unlikely to find goal, must approximate

# Minimax Search

- Core of many computer games
- Pertains primarily to:
  - Turn based games
  - Two players
  - Players with “perfect knowledge”

# Game tree (2-player, deterministic, turns)



# Game Tree

- Nodes are states
- Edges are decisions
- Levels are called “plys”

# Naïve Approach

- Given a game tree, what would be the most straightforward playing approach?
- Any potential problems?




# Minimax

- **Min**imizing the **max**imum possible loss
- Choose move which results in best state
  - Select highest expected score for you
- Assume opponent is playing optimally too
  - Will choose lowest expected score for you

# Minimax

- Perfect play for deterministic games
- Idea: choose move to position with highest **minimax value**  
= best achievable payoff against best play
- E.g., 2-ply game:



# Minimax algorithm

**function** MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

**return** the *action* in SUCCESSORS(*state*) with value *v*

---

**function** MAX-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

**return** *v*

---

**function** MIN-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

**return** *v*

# Properties of minimax

- Complete? Yes (if tree is finite)
- Optimal? Yes (against an optimal opponent)
- Time complexity?  $O(b^m)$
- Space complexity?  $O(bm)$  (depth-first exploration)
- For chess,  $b \approx 35$ ,  $m \approx 100$  for "reasonable" games  
→ exact solution completely infeasible

# Resource limits

Suppose we have 100 secs, explore  $10^4$  nodes/sec

→  $10^6$  nodes per move

Standard approach:

- **cutoff test:**  
e.g., depth limit (perhaps add **quiescence search**)
- **evaluation function**  
= estimated desirability of position

# Evaluation Functions

- Assign a utility score to a state
  - Different for players?
- Usually a range of integers
  - $[-1000, +1000]$
- $+\infty$  for win
- $-\infty$  for loss

# Cutting Off Search

- How to score a game before it ends?
  - You have to fudge it!
- Use a **heuristic** function to approximate state's utility

# Cutting off search

*MinimaxCutoff* is identical to *MinimaxValue* except

1. *Terminal?* is replaced by *Cutoff?*
2. *Utility* is replaced by *Eval*

Does it work in practice?

$$b^m = 10^6, b=35 \rightarrow m=4$$

4-ply lookahead is a hopeless chess player!

- 4-ply  $\approx$  human novice
- 8-ply  $\approx$  typical PC, human master
- 12-ply  $\approx$  Deep Blue, Kasparov

(A computer program which evaluates no further than its own legal moves plus the legal responses to those moves is searching to a depth of two-ply. )



# Example Evaluation Function

- For chess, typically **linear** weighted sum of **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g.,  $w_1 = 9$  with  
 $f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$ , etc.


# Evaluating States

- Assuming an ideal evaluation function, how would you make a move?
- Is this a good strategy with a bad function?

# Look Ahead


- Instead of only evaluating immediate future, look as far ahead as possible

# Look Ahead



# Bubbling Up

- Looking ahead allows utility values to “bubble up” to root of search tree



# Minimax Algorithm

- BESTMOVE function
- Inputs:
  - Board state
  - Depth bound
- Explores search tree to specified depth
- Output:
  - Best move

# Minimax Algorithm

```
BESTMOVE = proc (posn, depth)
  begin  (movelist, bestscore, bestm, try, tryscore) %local variables

    % posn: is the current BOARD CONFIGURATION FROM WHICH A MOVE
    %       MUST BE CHOSEN BY OUR INTELLIGENT AGENT COMPUTER

    % depth: MAXIMUM NUMBER OF PLIES TO LOOK AHEAD. THIS IS DETERMINED
    %       BY SPEED CONSTRAINTS AND COMPLEXITY OF THE GAME
    %       (i.e. branching factor b)

    % movelist: the list of all possible MOVES from the current posn
    % bestscore: the best "backed up" score found so far as we iterate
    %            thru the movelist
    % bestm:     the best move found so far that results in the bestscore
    % try:       just a tmp to hold returned values from recursive calls
    % tryscore:  just a tmp to hold returned scores from recursive calls
```

# Minimax Algorithm

```
if depth = 0 then return list(EVALUATE(posn), nil) fi;  
  %This ends the recursion at the bottom of the search space  
  %Note a two element list of values is returned.  
  
Movelist = POSSIBLE-MOVES(posn);  
  %According to the rules of the game, we generate all possible  
  %moves from the given board position.  
  
Bestscore = -infinity;  %initializations  
Bestm = nil;
```



# Minimax Algorithm

%We are now ready to scan the movelist and select the  
%best move. Note how we initialized Bestscore and Bestm.

```
while ( Movelist <> nil )
  repeat,
    try = BESTMOVE( NEWPOSITION(posn, first(Movelist)), depth-1);
    %Here is the main recursive call that expands and searches
    %the state space from the selected move.

    tryscore = - first(try); %recall BESTMOVE returns two values

    %Now we determine how well this current move did and whether
    %it should be selected as our best move found so far.

    if tryscore > Bestscore then
      do,
        Bestscore = tryscore;
        Bestm = first(Movelist);
      od;

    %Now we continue scanning down the list of moves to see
    %if we can find a better move than found so far.

    Movelist = rest(Movelist);
  taeper;
```

# Minimax Algorithm

```
%After scanning the entire Movelist, we have our best move so:
```

```
return( list(Bestscore, Bestm) );
```

```
nigeb; %End of procedure min-max.
```

# Minimax Algorithm

- Did you notice anything missing?

# Minimax Algorithm

- Did you notice anything missing?
- Where were Max-Value and Min-Value?

# Minimax Algorithm

- Did you notice anything missing?
- Where were Max-Value and Min-Value?
- What is going on here?

```
tryscore = - first(try); %recall BESTMOVE returns two values
```

# Minimax Algorithm

- Did you notice anything missing?
- Where were Max-Value and Min-Value?
- What is going on here?

```
tryscore = - first(try); %recall BESTMOVE returns two values
```

# Be Careful!

- Things to worry about?

# Complexity

- What is the space complexity of depth-bounded Minimax?



# Complexity

- What is the space complexity of depth-bounded Minimax?
  - Board size  $s$
  - Depth  $d$
  - Possible moves  $m$

# Complexity

- What is the space complexity of depth-bounded Minimax?
  - Board size  $s$
  - Depth  $d$
  - Possible moves  $m$
- $O(ds+m)$
- Board positions can be released as bubble up

# Minimax Algorithm

- Did I just do your project for you?

# Minimax Algorithm

- Did I just do your project for you?
- No!

# Minimax Algorithm

- Did I just do your project for you?
- No!
- You need to create:
  - Evaluation function
  - Move generator
  - did\_i\_win? function

# Isolation Clarification

- Standalone game clients
  - Opponent moves entered manually
  - Output your move on stdout
- Assignment is out of 100
- Tournament is single elimination
  - But there will be food!

# Recap

- What is a zero sum game?

# Recap

- What is a zero sum game?
- What is a game tree?



# Recap

- What is a zero sum game?
- What is a game tree?
- What is Minimax?

# Recap

- What is a zero sum game?
- What is a game tree?
- What is Minimax?
  - Why is it called that?

# Recap

- What is a zero sum game?
- What is a game tree?
- What is Minimax?
  - Why is it called that?
- What is its space complexity?

# Recap

- What is a zero sum game?
- What is a game tree?
- What is Minimax?
  - Why is it called that?
- What is its space complexity?
- How can the Minimax algorithm be simplified?

# Recap

- What is a zero sum game?
- What is a game tree?
- What is Minimax?
  - Why is it called that?
- What is its space complexity?
- How can the Minimax algorithm be simplified?
  - Will this work for all games?

# Recap

- What is a zero sum game?
- What is a game tree?
- What is Minimax?
  - Why is it called that?
- What is its space complexity?
- How can the Minimax algorithm be simplified?
  - Will this work for all games?

# Next Up

- Recall that minimax will produce optimal play against an optimal opponent if entire tree is searched
- Is the same true if a cutoff is used?


# Horizon Effect

- Your algorithm searches to depth  $n$
- What happens if:
  - Evaluation(s) at depth  $n$  is very positive
  - Evaluation(s) at depth  $n+1$  is very negative
- Or:
  - Evaluation(s) at depth  $n$  is very negative
  - Evaluation(s) at depth  $n+1$  is very positive
- Will this ever happen in practice?



# Local Maxima Problem

$$f(x,y) = e^{-(x^2+y^2)} + 2e^{-((x-1.7)^2+(y-1.7)^2)}$$



# Search Limitation Mitigation

- Sometimes it is useful to look deeper into game tree
- We could peak past the horizon...
- But how can you decide what nodes to explore?
  - Quiescence search

# Quiescence Search

- Human players have some intuition about move quality
  - “Interesting vs “boring”
  - “Promising” vs “dead end”
  - “Noisy” vs “quiet”
- Expand horizon for potential high impact moves
- Quiescence search adds this to Minimax

# Quiescence Search

- Additional search performed on leaf nodes
- if looks\_interesting(leaf\_node):  
    extend\_search\_depth(leaf\_node)  
else:  
    normal\_evaluation(leaf\_node)

# Quiescence Search

- What constitutes an “interesting” state?
  - Moves that substantially alter game state
  - Moves that cause large fluctuations in evaluation function output
- Chess example: capture moves
- Must be careful to prevent indefinite extension of search depth
  - Chess: checks vs captures


# Search Limitation Mitigation

- Do you always need to search the entire tree?
  - No!
- Sometimes it is useful to look *less deeply* into tree
- But how can you decide what branches to ignore?
  - Tree pruning

# Tree Pruning

- Moves chosen under assumption of optimal adversary
- You know the best move so far
- If you find a branch with a worse move, is there any point in looking further?
- Thought experiment: bag game

# Pruning Example





# Alpha-Beta Pruning


- During Minimax, keep track of two additional values
- Alpha
  - Your best score via any path
- Beta
  - Opponent's best score via any path

# Alpha-Beta Pruning

- Max player (you) will never make a move that could lead to a worse score for you
- Min player (opponent) will never make a move that could lead to a better score for you
- Stop evaluating a branch whenever:
  - A value greater than beta is found
  - A value less than alpha is found

# Why is it called $\alpha$ - $\beta$ ?

- $\alpha$  is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- If  $v$  is worse than  $\alpha$ , *max* will avoid it  
→ prune that branch
- Define  $\beta$  similarly for *min*



# Alpha-Beta Pruning

- Based on observation that for all viable paths utility value  $n$  will be  $\alpha \leq n \leq \beta$


# Alpha-Beta Pruning

- Initially,  $\alpha = -\text{infinity}$ ,  $\beta = \text{infinity}$




# Alpha-Beta Pruning

- As the search tree is traversed, the possible utility value window shrinks as
  - Alpha increases
  - Beta decreases



# Alpha-Beta Pruning

- Once there is no longer any overlap in the possible ranges of alpha and beta, it is safe to conclude that the current node is a dead end



# Minimax algorithm

**function** MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

**return** the *action* in SUCCESSORS(*state*) with value *v*

---

**function** MAX-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

**return** *v*

---

**function** MIN-VALUE(*state*) *returns a utility value*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

**for** *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

**return** *v*



# The $\alpha$ - $\beta$ algorithm

**function** ALPHA-BETA-SEARCH(*state*) *returns an action*

**inputs:** *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

**return** the *action* in SUCCESSORS(*state*) with value  $v$

---

**function** MAX-VALUE(*state*,  $\alpha$ ,  $\beta$ ) *returns a utility value*

**inputs:** *state*, current state in game

$\alpha$ , the value of the best alternative for MAX along the path to *state*

$\beta$ , the value of the best alternative for MIN along the path to *state*

**if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

**for**  $a, s$  in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$

$\alpha \leftarrow \text{MAX}(\alpha, v)$


**return**  $v$

# The $\alpha$ - $\beta$ algorithm


```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to state
            $\beta$ , the value of the best alternative for MIN along the path to state

  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```


# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example




# $\alpha$ - $\beta$ pruning example



# $\alpha$ - $\beta$ pruning example



# Another $\alpha$ - $\beta$ Pruning Example





# Minimax Psuedocode

```
BESTMOVE = proc (posn, depth)
  begin  (movelist, bestscore, bestm, try, tryscore) %local variables

    % posn: is the current BOARD CONFIGURATION FROM WHICH A MOVE
    %       MUST BE CHOSEN BY OUR INTELLIGENT AGENT COMPUTER

    % depth: MAXIMUM NUMBER OF PLIES TO LOOK AHEAD. THIS IS DETERMINED
    %       BY SPEED CONSTRAINTS AND COMPLEXITY OF THE GAME
    %       (i.e. branching factor b)

    % movelist: the list of all possible MOVES from the current posn
    % bestscore: the best "backed up" score found so far as we iterate
    %            thru the movelist
    % bestm:     the best move found so far that results in the bestscore
    % try:       just a tmp to hold returned values from recursive calls
    % tryscore:  just a tmp to hold returned scores from recursive calls
```

# Alpha-Beta Psuedocode

```
BESTMOVE = proc (Posn, Depth, Mybest, Herbest)
    %Note we added two additional parameters. Mybest should be
    %initialized to -infinity, while herbest is initialized to
    %+infinity when calling this version of BESTMOVE.

    begin    (Movelist, Bestscore, Bestm, Try, Tryscore) %local variables
```

# Minimax Psuedocode

```
if depth = 0 then return list(EVALUATE(posn), nil) fi;  
  %This ends the recursion at the bottom of the search space  
  %Note a two element list of values is returned.  
  
Movelist = POSSIBLE-MOVES(posn);  
  %According to the rules of the game, we generate all possible  
  %moves from the given board position.  
  
Bestscore = -infinity;  %initializations  
Bestm = nil;
```

# Alpha-Beta Psuedocode

```
if Depth = 0 then return list(EVALUATE(Posn), nil) fi;  
  %This ends the recursion at the bottom of the search space  
  %Note a two element list of values is returned.  
  
Movelist = POSSIBLE-MOVES(Posn);  
  %According to the rules of the game, we generate all possible  
  %moves from the given board position.  
  
Bestscore = Mybest  % note the change to initializations here  
Bestm = nil;
```

# Minimax Psuedocode

%We are now ready to scan the movelist and select the  
%best move. Note how we initialized Bestscore and Bestm.

```
while ( Movelist <> nil )  
  repeat,  
    try = BESTMOVE( NEWPOSITION(posn, first(Movelist)), depth-1);  
    %Here is the main recursive call that expands and searches  
    %the state space from the selected move.  
  
    tryscore = - first(try); %recall BESTMOVE returns two values  
  
    %Now we determine how well this current move did and whether  
    %it should be selected as our best move found so far.  
  
    if tryscore > Bestscore then  
      do,  
        Bestscore = tryscore;  
        Bestm = first(Movelist);  
      od;  
  
    %Now we continue scanning down the list of moves to see  
    %if we can find a better move than found so far.  
  
    Movelist = rest(Movelist);  
  taeper;
```

# Alpha-Beta Psuedocode

%We are now ready to scan the movelist and select the  
%best move. Note how we initialized Bestscore and Bestm.

```
while ( Movelist <> nil )
  repeat,
    Try =
      BESTMOVE( NEWPOSITION(posn, first(Movelist)),
                Depth-1,
                -Herbest,
                -Bestscore );
    %Here is the main recursive call that expands and searches
    %the state space from the selected move. But notice we
    %added the two additional parameters here. This makes
    %sense when we view the following conditional expressions.

    Tryscore = - first(Try); %recall BESTMOVE returns two values

    %Now we determine how well this current move did and whether
    %it should be selected as our best move found so far.

    if Tryscore > Bestscore then
      do,
        Bestscore = Tryscore;
        Bestm = first(Movelist);
      od;
```

# Minimax Algorithm

```
%After scanning the entire Movelist, we have our best move so:
```

```
return( list(Bestscore, Bestm) );
```

```
nigeb; %End of procedure min-max.
```

# Alpha-Beta Psuedocode

```
if Bestscore > Herbest then return ( list(Bestscore, Bestm) );
```

```
%Now we continue scanning down the list of moves to see  
%if we can find a better move then found so far.
```

```
    Movelist = rest(Movelist);  
taeper;
```

```
%After scanning the entire Movelist, we have our best move so:
```

```
return( list(Bestscore, Bestm) );
```

```
nigeb; %End of procedure alpha-beta.
```




# Tree Pruning vs Heuristics

- Search depth cut off may affect outcome of algorithm
- How about pruning?

# Move Ordering

- Does the order in which moves are listed have any impact of alpha-beta?



# Move Ordering

- Techniques for improving move ordering
- Apply evaluation function to nodes prior to expanding children
  - Search in descending order
  - But sacrifices search depth
- Cache results of previous algorithm

# Properties of $\alpha$ - $\beta$

- Pruning **does not** affect final result
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity =  $O(b^{m/2})$   
→ **doubles** depth of search
- A simple example of the value of reasoning about which computations are relevant (a form of **metareasoning**)

# Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a pre-computed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.

# Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a pre-computed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.

# Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a pre-computed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are too good.

# Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a pre-computed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are too good.
- Go: human champions refuse to compete against computers, who are too bad. In go,  $b > 300$ , so most programs use pattern knowledge bases to suggest plausible moves.



# Summary

- Games are fun to work on!
- They illustrate several important points about AI
- perfection is unattainable → must approximate
- good idea to think about what to think about